


Version 3.0 Specification

May 2006



Notice

© 2026 Microsoft Corporation. All rights reserved.

Microsoft, Windows, Visual Basic, Visual C#, and Visual C++ are either registered trademarks or trademarks of Microsoft Corporation in the U.S.A. and/or other countries/regions.

Other product and company names mentioned herein may be the trademarks of their respective owners.

Table of Contents

26. Overview of C# 3.0	1
26.1 Implicitly typed local variables	1
26.2 Extension methods.....	2
26.2.1 Declaring extension methods.....	2
26.2.2 Importing extension methods.....	3
26.2.3 Extension method invocations	3
26.3 Lambda expressions	4
26.3.1 Lambda expression conversions	6
26.3.2 Type inference	7
26.3.3 Overload resolution.....	8
26.4 Object and collection initializers	9
26.4.1 Object initializers	10
26.4.2 Collection initializers	12
26.5 Anonymous types	13
26.6 Implicitly typed arrays.....	14
26.7 Query expressions	15
26.7.1 Query expression translation.....	16
26.7.1.1 Select and groupby clauses with continuations	17
26.7.1.2 Explicit iteration variable types.....	17
26.7.1.3 Join clauses.....	18
26.7.1.4 Let and where clauses	19
26.7.1.5 Multiple generators	20
26.7.1.6 Orderby clauses	21
26.7.1.7 Select clauses.....	22
26.7.1.8 Groupby clauses	22
26.7.1.9 Transparent identifiers.....	22
26.7.2 The query expression pattern	24
26.8 Expression trees.....	25

26. Overview of C# 3.0

C# 3.0 (“C# Orcas”) introduces several language extensions that build on C# 2.0 to support the creation and use of higher order, functional style class libraries. The extensions enable construction of compositional APIs that have equal expressive power of query languages in domains such as relational databases and XML. The extensions include:

- Implicitly typed local variables, which permit the type of local variables to be inferred from the expressions used to initialize them.
- Extension methods, which make it possible to extend existing types and constructed types with additional methods.
- Lambda expressions, an evolution of anonymous methods that provides improved type inference and conversions to both delegate types and expression trees.
- Object initializers, which ease construction and initialization of objects.
- Anonymous types, which are tuple types automatically inferred and created from object initializers.
- Implicitly typed arrays, a form of array creation and initialization that infers the element type of the array from an array initializer.
- Query expressions, which provide a language integrated syntax for queries that is similar to relational and hierarchical query languages such as SQL and XQuery.
- Expression trees, which permit lambda expressions to be represented as data (expression trees) instead of as code (delegates).

This document is a technical overview of those features. The document makes reference to the C# Language Specification 1.2 (§1 through §18) and the C# Language Specification 2.0 (§19 through §25), both of which are available on the C# Language Home Page (<http://msdn.microsoft.com/vcsharp/language>).

26.1 Implicitly typed local variables

In an *implicitly typed local variable declaration*, the type of the local variable being declared is inferred from the expression used to initialize the variable. When a local variable declaration specifies `var` as the type and no type named `var` is in scope, the declaration is an implicitly typed local variable declaration. For example:

```
var i = 5;
var s = "Hello";
var d = 1.0;
var numbers = new int[] {1, 2, 3};
var orders = new Dictionary<int,Order>();
```

The implicitly typed local variable declarations above are precisely equivalent to the following explicitly typed declarations:

```
int i = 5;
string s = "Hello";
double d = 1.0;
int[] numbers = new int[] {1, 2, 3};
Dictionary<int,Order> orders = new Dictionary<int,Order>();
```

C# 3.0 SPECIFICATION

A local variable declarator in an implicitly typed local variable declaration is subject to the following restrictions:

- The declarator must include an initializer.
- The initializer must be an expression. The initializer cannot be an object or collection initializer (§26.4) by itself, but it can be a `new` expression that includes an object or collection initializer.
- The compile-time type of the initializer expression cannot be the null type.
- If the local variable declaration includes multiple declarators, the initializers must all have the same compile-time type.

The following are examples of incorrect implicitly typed local variable declarations:

```
var x;           // Error, no initializer to infer type from
var y = {1, 2, 3}; // Error, collection initializer not permitted
var z = null;    // Error, null type not permitted
```

For reasons of backward compatibility, when a local variable declaration specifies `var` as the type and a type named `var` is in scope, the declaration refers to that type; however, a warning is generated to call attention to the ambiguity. Since a type named `var` violates the established convention of starting type names with an upper case letter, this situation is unlikely to occur.

The *for-initializer* of a `for` statement (§8.8.3) and the *resource-acquisition* of a `using` statement (§8.13) can be an implicitly typed local variable declaration. Likewise, the iteration variable of a `foreach` statement (§8.8.4) may be declared as an implicitly typed local variable, in which case the type of the iteration variable is inferred to be the element type of the collection being enumerated. In the example

```
int[] numbers = { 1, 3, 5, 7, 9 };
foreach (var n in numbers) Console.WriteLine(n);
```

the type of `n` is inferred to be `int`, the element type of `numbers`.

26.2 Extension methods

Extension methods are static methods that can be invoked using instance method syntax. In effect, extension methods make it possible to extend existing types and constructed types with additional methods.

Note

Extension methods are less discoverable and more limited in functionality than instance methods. For those reasons, it is recommended that extension methods be used sparingly and only in situations where instance methods are not feasible or possible.

Extension members of other kinds, such as properties, events, and operators, are being considered but are currently not supported.

26.2.1 Declaring extension methods

Extension methods are declared by specifying the keyword `this` as a modifier on the first parameter of the methods. Extension methods can only be declared in static classes. The following is an example of a static class that declares two extension methods:

```
namespace Acme.Utilities
{
    public static class Extensions
    {
        public static int ToInt32(this string s) {
            return Int32.Parse(s);
        }
    }
}
```

```
public static T[] Slice<T>(this T[] source, int index, int count) {
    if (index < 0 || count < 0 || source.Length - index < count)
        throw new ArgumentException();
    T[] result = new T[count];
    Array.Copy(source, index, result, 0, count);
    return result;
}
}
```

Extension methods have all the capabilities of regular static methods. In addition, once imported, extension methods can be invoked using instance method syntax.

26.2.2 Importing extension methods

Extension methods are imported through *using-namespace-directives* (§9.3.2). In addition to importing the types contained in a namespace, a *using-namespace-directive* imports all extension methods in all static classes in the namespace. In effect, imported extension methods appear as additional methods on the types that are given by their first parameter and have lower precedence than regular instance methods. For example, when the `Acme.Utilities` namespace from the example above is imported with the *using-namespace-directive*

```
using Acme.Utilities;
```

it becomes possible to invoke the extension methods in the static class `Extensions` using instance method syntax:

```
string s = "1234";
int i = s.ToInt32();           // Same as Extensions.ToInt32(s)

int[] digits = {0, 1, 2, 3, 4, 5, 6, 7, 8, 9};
int[] a = digits.Slice(4, 3);  // Same as Extensions.Slice(digits, 4, 3)
```

26.2.3 Extension method invocations

The detailed rules for extension method invocation are described in the following. In a method invocation (§7.5.5.1) of one of the forms

```
expr . identifier ( )
expr . identifier ( args )
expr . identifier < typeargs > ( )
expr . identifier < typeargs > ( args )
```

if the normal processing of the invocation finds no applicable instance methods (specifically, if the set of candidate methods for the invocation is empty), an attempt is made to process the construct as an extension method invocation. The method invocation is first rewritten to one of the following, respectively:

```
identifier ( expr )
identifier ( expr , args )
identifier < typeargs > ( expr )
identifier < typeargs > ( expr , args )
```

The rewritten form is then processed as a static method invocation, except for the way in which *identifier* is resolved: Starting with the closest enclosing namespace declaration, continuing with each enclosing namespace declaration, and ending with the containing compilation unit, successive attempts are made to process the rewritten method invocation with a method group consisting of all accessible extension methods with the name

given by *identifier* imported by the namespace declaration's *using-namespace-directives*. The first method group that yields a non-empty set of candidate methods is the one chosen for the rewritten method invocation. If all attempts yield empty sets of candidate methods, a compile-time error occurs.

The preceding rules mean that instance methods take precedence over extension methods, and extension methods imported in inner namespace declarations take precedence over extension methods imported in outer namespace declarations. For example:

```
using N1;
namespace N1
{
    public static class E
    {
        public static void F(this object obj, int i) { }
        public static void F(this object obj, string s) { }
    }
}
class A { }
class B
{
    public void F(int i) { }
}
class C
{
    public void F(object obj) { }
}
class X
{
    static void Test(A a, B b, C c) {
        a.F(1);           // E.F(object, int)
        a.F("hello");     // E.F(object, string)

        b.F(1);           // B.F(int)
        b.F("hello");     // E.F(object, string)

        c.F(1);           // C.F(object)
        c.F("hello");     // C.F(object)
    }
}
```

In the example, B's method takes precedence over the first extension method, and C's method takes precedence over both extension methods.

26.3 Lambda expressions

C# 2.0 introduces anonymous methods, which allow code blocks to be written “in-line” where delegate values are expected. While anonymous methods provide much of the expressive power of functional programming languages, the anonymous method syntax is rather verbose and imperative in nature. ***Lambda expressions*** provide a more concise, functional syntax for writing anonymous methods.

A lambda expression is written as a parameter list, followed by the => token, followed by an expression or a statement block.

expression:
assignment
non-assignment-expression

non-assignment-expression:

conditional-expression

lambda-expression

query-expression

lambda-expression:

(lambda-parameter-list_{opt}) => lambda-expression-body

implicitly-typed-lambda-parameter => lambda-expression-body

lambda-parameter-list:

explicitly-typed-lambda-parameter-list

implicitly-typed-lambda-parameter-list

explicitly-typed-lambda-parameter-list

explicitly-typed-lambda-parameter

explicitly-typed-lambda-parameter-list , explicitly-typed-lambda-parameter

explicitly-typed-lambda-parameter:

parameter-modifier_{opt} type identifier

implicitly-typed-lambda-parameter-list

implicitly-typed-lambda-parameter

implicitly-typed-lambda-parameter-list , implicitly-typed-lambda-parameter

implicitly-typed-lambda-parameter:

identifier

lambda-expression-body:

expression

block

The parameters of a lambda expression can be explicitly or implicitly typed. In an explicitly typed parameter list, the type of each parameter is explicitly stated. In an implicitly typed parameter list, the types of the parameters are inferred from the context in which the lambda expression occurs—specifically, when the lambda expression is converted to a compatible delegate type, that delegate type provides the parameter types (§26.3.1).

In a lambda expression with a single, implicitly typed parameter, the parentheses may be omitted from the parameter list. In other words, a lambda expression of the form

(param) => expr

can be abbreviated to

param => expr

Some examples of lambda expressions follow below:

```
x => x + 1 // Implicitly typed, expression body
x => { return x + 1; } // Implicitly typed, statement body
(int x) => x + 1 // Explicitly typed, expression body
(int x) => { return x + 1; } // Explicitly typed, statement body
(x, y) => x * y // Multiple parameters
() => Console.WriteLine() // No parameters
```

C# 3.0 SPECIFICATION

In general, the specification of anonymous methods, provided in §21 of the C# 2.0 Specification, also applies to lambda expressions. Lambda expressions are a functional superset of anonymous methods, providing the following additional functionality:

- Lambda expressions permit parameter types to be omitted and inferred whereas anonymous methods require parameter types to be explicitly stated.
- The body of a lambda expression can be an expression or a statement block whereas the body of an anonymous method can only be a statement block.
- Lambda expressions passed as arguments participate in type argument inference (§26.3.2) and in method overload resolution (§26.3.3).
- Lambda expressions with an expression body can be converted to expression trees (§26.8).

26.3.1 Lambda expression conversions

Similar to an *anonymous-method-expression*, a *lambda-expression* is classified as a value with special conversion rules. The value does not have a type but can be implicitly converted to a compatible delegate type. Specifically, a delegate type *D* is compatible with a lambda-expression *L* provided:

- *D* and *L* have the same number of parameters.
- If *L* has an explicitly typed parameter list, each parameter in *D* has the same type and modifiers as the corresponding parameter in *L*.
- If *L* has an implicitly typed parameter list, *D* has no `ref` or `out` parameters.
- If *D* has a `void` return type and the body of *L* is an expression, when each parameter of *L* is given the type of the corresponding parameter in *D*, the body of *L* is a valid expression that would be permitted as a *statement-expression* (§8.6).
- If *D* has a `void` return type and the body of *L* is a statement block, when each parameter of *L* is given the type of the corresponding parameter in *D*, the body of *L* is a valid statement block in which no `return` statement specifies an expression.
- If *D* has a non-void return type and the body of *L* is an expression, when each parameter of *L* is given the type of the corresponding parameter in *D*, the body of *L* is a valid expression that is implicitly convertible to the return type of *D*.
- If *D* has a non-void return type and the body of *L* is a statement block, when each parameter of *L* is given the type of the corresponding parameter in *D*, the body of *L* is a valid statement block with a non-reachable end point in which each `return` statement specifies an expression that is implicitly convertible to the return type of *D*.

The examples that follow use a generic delegate type `Func<A, R>` which represents a function taking an argument of type *A* and returning a value of type *R*:

```
delegate R Func<A,R>(A arg);
```

In the assignments

```
Func<int,int> f1 = x => x + 1;           // Ok
Func<int,double> f2 = x => x + 1;        // Ok
Func<double,int> f3 = x => x + 1;        // Error
```

the parameter and return types of each lambda expression are determined from the type of the variable to which the lambda expression is assigned. The first assignment successfully converts the lambda expression to the delegate type `Func<int, int>` because, when `x` is given type `int`, `x + 1` is a valid expression that is implicitly convertible to type `int`. Likewise, the second assignment successfully converts the lambda expression to the delegate type `Func<int, double>` because the result of `x + 1` (of type `int`) is implicitly convertible to type `double`. However, the third assignment is a compile-time error because, when `x` is given type `double`, the result of `x + 1` (of type `double`) is not implicitly convertible to type `int`.

26.3.2 Type inference

When a generic method is called without specifying type arguments, a type inference process attempts to infer type arguments for the call. Lambda expressions passed as arguments to the generic method participate in this type inference process.

As described in §20.6.4, type inference first occurs independently for each argument. In this initial phase, nothing is inferred from arguments that are lambda expressions. However, following the initial phase, additional inferences are made from lambda expressions using an iterative process. Specifically, inferences are made as long as one or more arguments exist for which all of the following are true:

- The argument is a lambda expression, in the following called `L`, from which no inferences have yet been made.
- The corresponding parameter's type, in the following called `P`, is a delegate type with a return type that involves one or more method type parameters.
- `P` and `L` have the same number of parameters, and each parameter in `P` has the same modifiers as the corresponding parameter in `L`, or no modifiers if `L` has an implicitly typed parameter list.
- `P`'s parameter types involve no method type parameters or involve only method type parameters for which a consistent set of inferences have already been made.
- If `L` has an explicitly typed parameter list, when inferred types are substituted for method type parameters in `P`, each parameter in `P` has the same type as the corresponding parameter in `L`.
- If `L` has an implicitly typed parameter list, when inferred types are substituted for method type parameters in `P` and the resulting parameter types are given to the parameters of `L`, the body of `L` is a valid expression or statement block.
- A return type can be inferred for `L`, as described below.

For each such argument, inferences are made from that argument by relating the return type of `P` with the inferred return type of `L` and the new inferences are added to the accumulated set of inferences. This process is repeated until no further inferences can be made.

For purposes of type inference and overload resolution, the ***inferred return type*** of a lambda expression `L` is determined as follows:

- If the body of `L` is an expression, the type of that expression is the inferred return type of `L`.
- If the body of `L` is a statement block, if the set formed by the types of the expressions in the block's `return` statements contains exactly one type to which each type in the set is implicitly convertible, and if that type is not the null type, then that type is the inferred return type of `L`.
- Otherwise, a return type cannot be inferred for `L`.

C# 3.0 SPECIFICATION

As an example of type inference involving lambda expressions, consider the `Select` extension method declared in the `System.Query.Sequence` class:

```
namespace System.Query
{
    public static class Sequence
    {
        public static IEnumerable<S> Select<T,S>(
            this IEnumerable<T> source,
            Func<T,S> selector)
        {
            foreach (T element in source) yield return selector(element);
        }
    }
}
```

Assuming the `System.Query` namespace was imported with a `using` clause, and given a class `Customer` with a `Name` property of type `string`, the `Select` method can be used to select the names of a list of customers:

```
List<Customer> customers = GetCustomerList();
IEnumerable<string> names = customers.Select(c => c.Name);
```

The extension method invocation (§26.2.3) of `Select` is processed by rewriting the invocation to a static method invocation:

```
IEnumerable<string> names = Sequence.Select(customers, c => c.Name);
```

Since type arguments were not explicitly specified, type inference is used to infer the type arguments. First, the `customers` argument is related to the `source` parameter, inferring `T` to be `Customer`. Then, using the lambda expression type inference process described above, `c` is given type `Customer`, and the expression `c.Name` is related to the return type of the `selector` parameter, inferring `S` to be `string`. Thus, the invocation is equivalent to

```
Sequence.Select<Customer,string>(customers, (Customer c) => c.Name)
```

and the result is of type `IEnumerable<string>`.

The following example demonstrates how lambda expression type inference allows type information to “flow” between arguments in a generic method invocation. Given the method

```
static Z F<X,Y,Z>(X value, Func<X,Y> f1, Func<Y,Z> f2) {
    return f2(f1(value));
}
```

type inference for the invocation

```
double seconds = F("1:15:30", s => TimeSpan.Parse(s), t => t.TotalSeconds);
```

proceeds as follows: First, the argument `"1:15:30"` is related to the `value` parameter, inferring `X` to be `string`. Then, the parameter of the first lambda expression, `s`, is given the inferred type `string`, and the expression `TimeSpan.Parse(s)` is related to the return type of `f1`, inferring `Y` to be `System.TimeSpan`. Finally, the parameter of the second lambda expression, `t`, is given the inferred type `System.TimeSpan`, and the expression `t.TotalSeconds` is related to the return type of `f2`, inferring `Z` to be `double`. Thus, the result of the invocation is of type `double`.

26.3.3 Overload resolution

Lambda expressions in an argument list affect overload resolution in certain situations.

The following rule augments §7.4.2.3: Given a lambda expression `L` for which an inferred return type (§26.3.2) exists, an implicit conversion of `L` to a delegate type `D1` is a better conversion than an implicit conversion of `L` to

a delegate type D_2 if D_1 and D_2 have identical parameter lists and the implicit conversion from L 's inferred return type to D_1 's return type is a better conversion than the implicit conversion from L 's inferred return type to D_2 's return type. If these conditions are not true, neither conversion is better.

The following example illustrates the effect of this rule.

```
class ItemList<T>: List<T>
{
    public int Sum<T>(Func<T,int> selector) {
        int sum = 0;
        foreach (T item in this) sum += selector(item);
        return sum;
    }

    public double Sum<T>(Func<T,double> selector) {
        double sum = 0;
        foreach (T item in this) sum += selector(item);
        return sum;
    }
}
```

The `ItemList<T>` class has two `Sum` methods. Each takes a `selector` argument, which extracts the value to sum over from a list item. The extracted value can be either an `int` or a `double` and the resulting sum is likewise either an `int` or a `double`.

The `Sum` methods could for example be used to compute sums from a list of detail lines in an order.

```
class Detail
{
    public int UnitCount;
    public double UnitPrice;
    ...
}

void ComputeSums() {
    ItemList<Detail> orderDetails = GetOrderDetails(...);
    int totalUnits = orderDetails.Sum(d => d.UnitCount);
    double orderTotal = orderDetails.Sum(d => d.UnitPrice * d.UnitCount);
    ...
}
```

In the first invocation of `orderDetails.Sum`, both `Sum` methods are applicable because the lambda expression `d => d.UnitCount` is compatible with both `Func<Detail,int>` and `Func<Detail,double>`. However, overload resolution picks the first `Sum` method because the conversion to `Func<Detail,int>` is better than the conversion to `Func<Detail,double>`.

In the second invocation of `orderDetails.Sum`, only the second `Sum` method is applicable because the lambda expression `d => d.UnitPrice * d.UnitCount` produces a value of type `double`. Thus, overload resolution picks the second `Sum` method for that invocation.

26.4 Object and collection initializers

An object creation expression (§7.5.10.1) may include an object or collection initializer which initializes the members of the newly created object or the elements of the newly created collection.

object-creation-expression:

```
new type ( argument-listopt ) object-or-collection-initializeropt
new type object-or-collection-initializer
```

object-or-collection-initializer:
object-initializer
collection-initializer

An object creation expression can omit the constructor argument list and enclosing parentheses provided it includes an object or collection initializer. Omitting the constructor argument list and enclosing parentheses is equivalent to specifying an empty argument list.

Execution of an object creation expression that includes an object or collection initializer consists of first invoking the instance constructor and then performing the member or element initializations specified by the object or collection initializer.

It is not possible for an object or collection initializer to refer to the object instance being initialized.

26.4.1 Object initializers

An object initializer specifies values for one or more fields or properties of an object.

object-initializer:
 { *member-initializer-list*_{opt} }
 { *member-initializer-list* , }

member-initializer-list:
member-initializer
member-initializer-list , *member-initializer*

member-initializer:
identifier = *initializer-value*

initializer-value:
expression
object-or-collection-initializer

An object initializer consists of a sequence of member initializers, enclosed by { and } tokens and separated by commas. Each member initializer must name an accessible field or property of the object being initialized, followed by an equals sign and an expression or an object or collection initializer. It is an error for an object initializer to include more than one member initializer for the same field or property.

A member initializer that specifies an expression after the equals sign is processed in the same way as an assignment (§7.13.1) to the field or property.

A member initializer that specifies an object initializer after the equals sign is an initialization of an embedded object. Instead of assigning a new value to the field or property, the assignments in the object initializer are treated as assignments to members of the field or property. A property of a value type cannot be initialized using this construct.

A member initializer that specifies a collection initializer after the equals sign is an initialization of an embedded collection. Instead of assigning a new collection to the field or property, the elements given in the initializer are added to the collection referenced by the field or property. The field or property must be of a collection type that satisfies the requirements specified in §26.4.2.

The following class represents a point with two coordinates:

```
public class Point
{
    int x, y;
```

```
        public int X { get { return x; } set { x = value; } }
        public int Y { get { return y; } set { y = value; } }
    }
```

An instance of `Point` can be created and initialized as follows:

```
var a = new Point { X = 0, Y = 1 };
```

which has the same effect as

```
var a = new Point();
a.X = 0;
a.Y = 1;
```

The following class represents a rectangle created from two points:

```
public class Rectangle
{
    Point p1, p2;

    public Point P1 { get { return p1; } set { p1 = value; } }
    public Point P2 { get { return p2; } set { p2 = value; } }
}
```

An instance of `Rectangle` can be created and initialized as follows:

```
var r = new Rectangle {
    P1 = new Point { X = 0, Y = 1 },
    P2 = new Point { X = 2, Y = 3 }
};
```

which has the same effect as

```
var r = new Rectangle();
var __p1 = new Point();
__p1.X = 0;
__p1.Y = 1;
r.P1 = __p1;
var __p2 = new Point();
__p2.X = 2;
__p2.Y = 3;
r.P2 = __p2;
```

where `__p1` and `__p2` are temporary variables that are otherwise invisible and inaccessible.

If `Rectangle`'s constructor allocates the two embedded `Point` instances

```
public class Rectangle
{
    Point p1 = new Point();
    Point p2 = new Point();

    public Point P1 { get { return p1; } }
    public Point P2 { get { return p2; } }
}
```

the following construct can be used to initialize the embedded `Point` instances instead of assigning new instances:

```
var r = new Rectangle {
    P1 = { X = 0, Y = 1 },
    P2 = { X = 2, Y = 3 }
};
```

which has the same effect as

```
var r = new Rectangle();
r.P1.X = 0;
r.P1.Y = 1;
r.P2.X = 2;
r.P2.Y = 3;
```

26.4.2 Collection initializers

A collection initializer specifies the elements of a collection.

```
collection-initializer:
    { element-initializer-list }
    { element-initializer-list , }

element-initializer-list:
    element-initializer
    element-initializer-list , element-initializer

element-initializer:
    non-assignment-expression
```

A collection initializer consists of a sequence of element initializers, enclosed by { and } tokens and separated by commas. Each element initializer specifies an element to be added to the collection object being initialized. To avoid ambiguity with member initializers, element initializers cannot be assignment expressions. The *non-assignment-expression* production is defined in §26.3.

The following is an example of an object creation expression that includes a collection initializer:

```
List<int> digits = new List<int> { 0, 1, 2, 3, 4, 5, 6, 7, 8, 9 };
```

The collection object to which a collection initializer is applied must be of a type that implements `System.Collections.Generic.ICollection<T>` for exactly one `T`. Furthermore, an implicit conversion (§6.1) must exist from the type of each element initializer to `T`. A compile-time error occurs if these requirements are not satisfied. A collection initializer invokes the `ICollection<T>.Add(T)` method for each specified element in order.

The following class represents a contact with a name and a list of phone numbers:

```
public class Contact
{
    string name;
    List<string> phoneNumbers = new List<string>();

    public string Name { get { return name; } set { name = value; } }
    public List<string> PhoneNumbers { get { return phoneNumbers; } }
}
```

A `List<Contact>` can be created and initialized as follows:

```
var contacts = new List<Contact> {
    new Contact {
        Name = "Chris Smith",
        PhoneNumbers = { "206-555-0101", "425-882-8080" }
    },
    new Contact {
        Name = "Bob Harris",
        PhoneNumbers = { "650-555-0199" }
    }
};
```

which has the same effect as

```
var contacts = new List<Contact>();
var __c1 = new Contact();
__c1.Name = "Chris Smith";
__c1.PhoneNumbers.Add("206-555-0101");
__c1.PhoneNumbers.Add("425-882-8080");
contacts.Add(__c1);
var __c2 = new Contact();
__c2.Name = "Bob Harris";
__c2.PhoneNumbers.Add("650-555-0199");
contacts.Add(__c2);
```

where __c1 and __c2 are temporary variables that are otherwise invisible and inaccessible.

26.5 Anonymous types

C# 3.0 permits the `new` operator to be used with an anonymous object initializer to create an object of an anonymous type.

primary-no-array-creation-expression:

*...
anonymous-object-creation-expression*

anonymous-object-creation-expression:

new anonymous-object-initializer

anonymous-object-initializer:

*{ member-declarator-list_{opt} }
{ member-declarator-list , }*

member-declarator-list:

*member-declarator
member-declarator-list , member-declarator*

member-declarator:

*simple-name
member-access
identifier = expression*

An anonymous object initializer declares an anonymous type and returns an instance of that type. An anonymous type is a nameless class type that inherits directly from `object`. The members of an anonymous type are a sequence of read/write properties inferred from the object initializer(s) used to create instances of the type. Specifically, an anonymous object initializer of the form

```
new { p1 = e1 , p2 = e2 , ... pn = en }
```

declares an anonymous type of the form

```
class __Anonymous1
{
    private T1 f1 ;
    private T2 f2 ;
    ...
    private Tn fn ;

    public T1 p1 { get { return f1 ; } set { f1 = value ; } }
    public T2 p2 { get { return f2 ; } set { f2 = value ; } }
    ...
    public T1 p1 { get { return f1 ; } set { f1 = value ; } }
}
```

C# 3.0 SPECIFICATION

where each T_x is the type of the corresponding expression e_x . It is a compile-time error for an expression in an anonymous object initializer to be of the null type.

The name of an anonymous type is automatically generated by the compiler and cannot be referenced in program text.

Within the same program, two anonymous object initializers that specify a sequence of properties of the same names and types in the same order will produce instances of the same anonymous type. (This definition includes the order of the properties because it is observable and material in certain circumstances, such as reflection.)

In the example

```
var p1 = new { Name = "Lawnmower", Price = 495.00 };
var p2 = new { Name = "Shovel", Price = 26.95 };
p1 = p2;
```

the assignment on the last line is permitted because `p1` and `p2` are of the same anonymous type.

A member declarator can be abbreviated to a simple name (§7.5.2) or a member access (§7.5.4). This is called a ***projection initializer*** and is shorthand for a declaration of and assignment to a property with the same name. Specifically, member declarators of the forms

$$identif\grave{e}r \qquad \qquad \qquad expr \ . \ identif\grave{e}r$$

are precisely equivalent to the following, respectively:

$$\textit{identifer} = \textit{identifer} \qquad \textit{identifer} = \textit{expr} . \textit{identifer}$$

Thus, in a projection initializer the *identifier* selects both the value and the field or property to which the value is assigned. Intuitively, a projection initializer projects not just a value, but also the name of the value.

26.6 Implicitly typed arrays

The syntax of array creation expressions (§7.5.10.2) is extended to support implicitly typed array creation expressions:

```
array-creation-expression:
    ...
    new [ ] array-initializer
```

In an implicitly typed array creation expression, the type of the array instance is inferred from the elements specified in the array initializer. Specifically, the set formed by the types of the expressions in the array initializer must contain exactly one type to which each type in the set is implicitly convertible, and if that type is not the null type, an array of that type is created. If exactly one type cannot be inferred, or if the inferred type is the null type, a compile-time error occurs.

The following are examples of implicitly typed array creation expressions:

```
var a = new[] { 1, 10, 100, 1000 };           // int[]
var b = new[] { 1, 1.5, 2, 2.5 };             // double[]
var c = new[] { "hello", null, "world" };     // string[]
var d = new[] { 1, "one", 2, "two" };         // Error
```

The last expression causes a compile-time error because neither `int` nor `string` is implicitly convertible to the other. An explicitly typed array creation expression must be used in this case, for example specifying the type to be `object[]`. Alternatively, one of the elements can be cast to a common base type, which would then become the inferred element type.

Implicitly typed array creation expressions can be combined with anonymous object initializers to create anonymously typed data structures. For example:

```
var contacts = new[] {  
    new {  
        Name = "Chris Smith",  
        PhoneNumbers = new[] { "206-555-0101", "425-882-8080" }  
    },  
    new {  
        Name = "Bob Harris",  
        PhoneNumbers = new[] { "650-555-0199" }  
    }  
};
```

26.7 Query expressions

Query expressions provide a language integrated syntax for queries that is similar to relational and hierarchical query languages such as SQL and XQuery.

query-expression:

from-clause query-body

from-clause:

from type_{opt} identifier in expression join-clauses_{opt}

join-clauses:

join-clause

join-clauses join-clause

join-clause:

join type_{opt} identifier in expression on expression equals expression

join type_{opt} identifier in expression on expression equals expression into identifier

query-body:

from-let-where-clauses_{opt} orderby-clause_{opt} select-or-group-clause query-continuation_{opt}

from-let-where-clauses:

from-let-where-clause

from-let-where-clauses from-let-where-clause

from-let-where-clause:

from-clause

let-clause

where-clause

let-clause:

let identifier = expression

where-clause:

where boolean-expression

orderby-clause:

orderby orderings

orderings:

ordering

orderings , ordering

C# 3.0 SPECIFICATION

ordering:
 expression ordering-direction_{opt}

ordering-direction:
 ascending
 descending

select-or-group-clause:
 select-clause
 group-clause

select-clause:
 select *expression*

group-clause:
 group *expression* by *expression*

query-continuation:
 into *identifier join-clauses_{opt} query-body*

A *query-expression* is classified as a *non-assignment-expression*, the definition of which occurs in §26.3.

A query expression begins with a **from** clause and ends with either a **select** or **group** clause. The initial **from** clause can be followed by zero or more **from**, **let**, or **where** clauses. Each **from** clause is a generator that introduces one or more iteration variables ranging over a sequence or a join of multiple sequences, each **let** clause computes a value and introduces an identifier representing that value, and each **where** clause is a filter that excludes items from the result. The final **select** or **group** clause specifies the shape of the result in terms of the iteration variable(s). The **select** or **group** clause may be preceded by an **orderby** clause that specifies an ordering for the result. Finally, an **into** clause can be used to “splice” queries by treating the results of one query as a generator in a subsequent query.

26.7.1 Query expression translation

The C# 3.0 language does not specify the exact execution semantics of query expressions. Rather, C# 3.0 translates query expressions into invocations of methods that adhere to the *query expression pattern*. Specifically, query expressions are translated into invocations of methods named **where**, **Select**, **SelectMany**, **Join**, **GroupJoin**, **OrderBy**, **OrderByDescending**, **ThenBy**, **ThenByDescending**, **GroupBy**, and **Cast** that are expected to have particular signatures and result types, as described in §26.7.2. These methods can be instance methods of the object being queried or extension methods that are external to the object, and they implement the actual execution of the query.

The translation from query expressions to method invocations is a syntactic mapping that occurs before any type binding or overload resolution has been performed. The translation is guaranteed to be syntactically correct, but it is not guaranteed to produce semantically correct C# code. Following translation of query expressions, the resulting method invocations are processed as regular method invocations, and this may in turn uncover errors, for example if the methods do not exist, if arguments have wrong types, or if the methods are generic and type inference fails.

A query expression is processed by repeatedly applying the following translations until no further reductions are possible. The translations are listed in order of precedence: each section assumes that the translations in the preceding sections have been performed exhaustively.

Certain translations inject iteration variables with *transparent identifiers* denoted by *. The special properties of transparent identifiers are discussed further in §26.7.1.9.

26.7.1.1 Select and groupby clauses with continuations

A query expression with a continuation

```
from ... into x ...
```

is translated into

```
from x in ( from ... ) ...
```

The translations in the following sections assume that queries have no `into` continuations.

The example

```
from c in customers
group c by c.Country into g
select new { Country = g.Key, CustCount = g.Count() }
```

is translated into

```
from g in
    from c in customers
    group c by c.Country
select new { Country = g.Key, CustCount = g.Count() }
```

the final translation of which is

```
customers.
GroupBy(c => c.Country).
Select(g => new { Country = g.Key, CustCount = g.Count() })
```

26.7.1.2 Explicit iteration variable types

Note

The C# 3.0 compiler included in the May 2006 Technology Preview does not support explicitly typed iteration variables. Explicit typing can be provided by manually invoking the `Cast<T>` operator.

A `from` clause that explicitly specifies an iteration variable type

```
from T x in e
```

is translated into

```
from x in ( e ) . Cast < T > ( )
```

A `join` clause that explicitly specifies an iteration variable type

```
join T x in e on k1 equals k2
```

is translated into

```
join x in ( e ) . Cast < T > ( ) on k1 equals k2
```

The translations in the following sections assume that queries have no explicit iteration variable types.

The example

```
from Customer c in customers
where c.City == "London"
select c
```

is translated into

C# 3.0 SPECIFICATION

```
from c in customers.Cast<Customer>()
where c.City == "London"
select c
```

the final translation of which is

```
customers.
Cast<Customer>().
Where(c => c.City == "London")
```

Explicit iteration variable types are useful for querying collections that implement the non-generic `IEnumerable` interface, but not the generic `IEnumerable<T>` interface. In the example above, this would be the case if `customers` were of type `ArrayList`.

26.7.1.3 Join clauses

A `from` clause followed by a `join` clause without an `into` followed by a `select` clause

```
from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$  select  $v$ 
```

is translated into

```
from  $t$  in (  $e_1$  ) . Join (  $e_2$  ,  $x_1 \Rightarrow k_1$  ,  $x_2 \Rightarrow k_2$  , (  $x_1$  ,  $x_2$  )  $\Rightarrow v$  )
select  $t$ 
```

where t is a compiler-generated unique identifier.

A `from` clause followed by a `join` clause with an `into` followed by a `select` clause

```
from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$  into  $g$  select  $v$ 
```

is translated into

```
from  $t$  in (  $e_1$  ) . GroupJoin (  $e_2$  ,  $x_1 \Rightarrow k_1$  ,  $x_2 \Rightarrow k_2$  , (  $x_1$  ,  $g$  )  $\Rightarrow v$  )
select  $t$ 
```

where t is a compiler-generated unique identifier.

A `from` clause followed by a `join` clause without an `into` followed by something other than a `select` clause

```
from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$ 
```

is translated into

```
from * in (
    from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$  select new {  $x_1$  ,  $x_2$  }
)
```

A `from` clause followed by a `join` clause with an `into` followed by something other than a `select` clause

```
from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$  into  $g$ 
```

is translated into

```
from * in (
    from  $x_1$  in  $e_1$  join  $x_2$  in  $e_2$  on  $k_1$  equals  $k_2$  into  $g$  select new {  $x_1$  ,  $g$  }
)
```

The translations in the following sections assume that queries have no join clauses.

The example

```
from c in customers
join o in orders on c.CustomerID equals o.CustomerID
select new { c.Name, o.OrderDate, o.Total }
```

has the final translation

```
customers.Join(orders, c => c.CustomerID, o => o.CustomerID,
(c, o) => new { c.Name, o.OrderDate, o.Total })
```

The example

```
from c in customers
join o in orders on c.CustomerID equals o.CustomerID into co
let n = co.Count()
where n >= 10
select new { c.Name, OrderCount = n }
```

is translated into

```
from * in
  from c in customers
  join o in orders on c.CustomerID equals o.CustomerID into co
  select new { c, co }
let n = co.Count()
where n >= 10
select new { c.Name, OrderCount = n }
```

the final translation of which is

```
customers.
GroupJoin(orders, c => c.CustomerID, o => o.CustomerID,
(c, co) => new { c, co }).
Select(x => new { x, n = x.co.Count() }).
Where(y => y.n >= 10).
Select(y => new { y.x.c.Name, OrderCount = y.n})
```

where x and y are compiler generated identifiers that are otherwise invisible and inaccessible.

26.7.1.4 Let and where clauses

A from clause immediately followed by a let clause

```
from x in e let y = f
```

is translated into

```
from * in (
  from x in e select new { x , y = f }
)
```

A from clause immediately followed by a where clause

```
from x in e where f
```

is translated into

```
from x in ( e ) . where ( x => f )
```

The translations in the following sections assume that queries have no let or where clauses.

C# 3.0 SPECIFICATION

The example

```
from o in orders
let t = o.Details.Sum(d => d.UnitPrice * d.Quantity)
where t >= 1000
select new { o.OrderID, Total = t }
```

is translated into

```
from * in
    from o in orders
    select new { o, t = o.Details.Sum(d => d.UnitPrice * d.Quantity) }
where t >= 1000
select new { o.OrderID, Total = t }
```

the final translation of which is

```
orders.
Select(o => new { o, t = o.Details.Sum(d => d.UnitPrice * d.Quantity) }).
Where(x => x.t >= 1000).
Select(x => new { x.o.OrderID, Total = x.t })
```

where x is a compiler generated identifier that is otherwise invisible and inaccessible.

26.7.1.5 Multiple generators

Two from clauses followed by a select clause

```
from x1 in e1 from x2 in e2 select v
```

is translated into

```
from t in ( e1 ) . SelectMany ( x1 => from x2 in e2 select v ) select t
```

where t is a compiler-generated unique identifier.

Two or more from clauses followed by something other than a select clause

```
from x1 in e1 from x2 in e2 ...
```

is translated into

```
from * in (
    from x1 in e1 from x2 in e2 ... select new { x1 , x2 ... }
)
```

The translations in the following sections assume that queries have only one from clause.

The example

```
from c in customers
from o in c.Orders
select new { c.Name, o.OrderID, o.Total }
```

is translated into

```
from t in
    customers.SelectMany(c =>
        from o in c.Orders
        select new { c.Name, o.OrderID, o.Total }
    )
select t
```

the final translation of which is

```
customers.  
SelectMany(c =>  
    c.Orders.  
        Select(o => new { c.Name, o.OrderID, o.Total } })  
)
```

The example

```
from c in customers  
from o in c.Orders  
orderby o.Total descending  
select new { c.Name, o.OrderID, o.Total }
```

is translated into

```
from * in  
    from c in customers  
    from o in c.Orders  
    select new { c, o }  
orderby o.Total descending  
select new { c.Name, o.OrderID, o.Total }
```

the final translation of which is

```
customers.  
SelectMany(c =>  
    c.Orders.  
        Select(o => new { c, o } )  
    ).  
OrderByDescending(x => x.o.Total).  
Select(x => new { x.c.Name, x.o.OrderID, x.o.Total } )
```

where x is a compiler generated identifier that is otherwise invisible and inaccessible.

26.7.1.6 Orderby clauses

A query expression having an orderby clause

```
from x in e orderby  $k_1$  ,  $k_2$  ...
```

is translated into

```
from x in ( e ) . OrderBy ( x =>  $k_1$  ) . ThenBy ( x =>  $k_2$  ) ...
```

If an ordering clause specifies a descending direction indicator, an invocation of `OrderByDescending` or `ThenByDescending` is produced instead.

The translations in the following sections assume that queries have no orderby clause.

The example

```
from o in orders  
orderby o.Customer.Name, o.Total descending  
select o
```

has the final translation

```
orders.  
OrderBy(o => o.Customer.Name).  
ThenByDescending(o => o.Total)
```

26.7.1.7 Select clauses

A query expression of the form

```
from x in e select v
```

is translated into

```
( e ) . Select ( x => v )
```

except when v is the identifier x , the translation is simply

```
( e )
```

For example

```
from c in customers  
select c
```

is simply translated into

```
customers
```

26.7.1.8 Groupby clauses

A query expression of the form

```
from x in e group v by k
```

is translated into

```
( e ) . GroupBy ( x => k , x => v )
```

except when v is the identifier x , the translation is

```
( e ) . GroupBy ( x => k )
```

The example

```
from c in customers  
group c.Name by c.Country
```

is translated into

```
customers.  
GroupBy(c => c.Country, c => c.Name)
```

26.7.1.9 Transparent identifiers

Certain translations inject iteration variables with *transparent identifiers* denoted by *. Transparent identifiers are not a proper language feature; they exist only as an intermediate step in the query expression translation process.

When a query translation injects a transparent identifier, further translation steps propagate the transparent identifier into lambda expressions and anonymous object initializers. In those contexts, transparent identifiers have the following behavior:

- When a transparent identifier occurs as a parameter in a lambda expression, the members of the associated anonymous type are automatically in scope in the body of the lambda expression.
- When a member with a transparent identifier is in scope, its members are in scope as well.

- When a transparent identifier occurs as a member declarator in an anonymous object initializer, it introduces a member with a transparent identifier.

The example

```
from c in customers
from o in c.Orders
orderby o.Total descending
select new { c.Name, o.Total }
```

is translated into

```
from * in
    from c in customers
    from o in c.Orders
    select new { c, o }
orderby o.Total descending
select new { c.Name, o.Total }
```

which is further translated into

```
customers.
SelectMany(c => c.Orders.Select(o => new { c, o })).
OrderByDescending(* => o.Total).
Select(* => new { c.Name, o.Total })
```

which, when transparent identifiers are erased, is equivalent to

```
customers.
SelectMany(c => c.Orders.Select(o => new { c, o })).
OrderByDescending(x => x.o.Total).
Select(x => new { x.c.Name, x.o.Total })
```

where x is a compiler generated identifier that is otherwise invisible and inaccessible.

The example

```
from c in customers
join o in orders on c.CustomerID equals o.CustomerID
join d in details on o.OrderID equals d.OrderID
join p in products on d.ProductID equals p.ProductID
select new { c.Name, o.OrderDate, p.ProductName }
```

is translated into

```
from * in
    from * in
        from * in
            from c in customers
            join o in orders on c.CustomerID equals o.CustomerID
            select new { c, o }
        join d in details on o.OrderID equals d.OrderID
        select new { *, d }
    join p in products on d.ProductID equals p.ProductID
    select new { *, p }
select new { c.Name, o.OrderDate, p.ProductName }
```

which is further reduced to

```

customers.
Join(orders, c => c.CustomerID, o => o.CustomerID,
    (c, o) => new { c, o }).
Join(details, * => o.OrderID, d => d.OrderID,
    (*, d) => new { *, d }).
Join(products, * => d.ProductID, p => p.ProductID,
    (*, p) => new { *, p }).
Select(* => new { c.Name, o.OrderDate, p.ProductName })

```

the final translation of which is

```

customers.
Join(orders, c => c.CustomerID, o => o.CustomerID,
    (c, o) => new { c, o }).
Join(details, x => x.o.OrderID, d => d.OrderID,
    (x, d) => new { x, d }).
Join(products, y => y.d.ProductID, p => p.ProductID,
    (y, p) => new { y, p }).
Select(z => new { z.y.x.c.Name, z.y.x.o.OrderDate, z.p.ProductName })

```

where x, y, and z are compiler generated identifiers that are otherwise invisible and inaccessible.

26.7.2 The query expression pattern

The *Query Expression Pattern* establishes a pattern of methods that types can implement to support query expressions. Because query expressions are translated to method invocations by means of a syntactic mapping, types have considerable flexibility in how they implement the query expression pattern. For example, the methods of the pattern can be implemented as instance methods or as extension methods because the two have the same invocation syntax, and the methods can request delegates or expression trees because lambda expressions are convertible to both.

The recommended shape of a generic type `C<T>` that supports the query expression pattern is shown below. A generic type is used in order to illustrate the proper relationships between parameter and result types, but it is possible to implement the pattern for non-generic types as well.

```

delegate R Func<T1,R>(T1 arg1);
delegate R Func<T1,T2,R>(T1 arg1, T2 arg2);
class C
{
    public C<T> Cast<T>();
}
class C<T>
{
    public C<T> Where(Func<T,bool> predicate);
    public C<U> Select<U>(Func<T,U> selector);
    public C<U> SelectMany<U>(Func<T,C<U>> selector);
    public C<V> Join<U,K,V>(C<U> inner, Func<T,K> outerKeySelector,
        Func<U,K> innerKeySelector, Func<T,U,V> resultSelector);
    public C<V> GroupJoin<U,K,V>(C<U> inner, Func<T,K> outerKeySelector,
        Func<U,K> innerKeySelector, Func<T,C<U>,V> resultSelector);
    public O<T> OrderBy<K>(Func<T,K> keySelector);
    public O<T> OrderByDescending<K>(Func<T,K> keySelector);
    public C<G<K,T>> GroupBy<K>(Func<T,K> keySelector);
}

```

```
public C<G<K,E>> GroupBy<K,E>(Func<T,K> keySelector,
    Func<T,E> elementSelector);
}

class O<T> : C<T>
{
    public O<T> ThenBy<K>(Func<T,K> keySelector);
    public O<T> ThenByDescending<K>(Func<T,K> keySelector);
}

class G<K,T> : C<T>
{
    public K Key { get; }
}
```

The methods above use the generic delegate types `Func<T1, R>` and `Func<T1, T2, R>`, but they could equally well have used other delegate or expression tree types with the same relationships in parameter and result types.

Notice the recommended relationship between `C<T>` and `O<T>` which ensures that the `ThenBy` and `ThenByDescending` methods are available only on the result of an `OrderBy` or `OrderByDescending`. Also notice the recommended shape of the result of `GroupBy`—a sequence of sequences, where each inner sequence has an additional `Key` property.

The Standard Query Operators (described in a separate specification) provide an implementation of the query operator pattern for any type that implements the `System.Collections.Generic.IEnumerable<T>` interface.

26.8 Expression trees

Expression trees permit lambda expressions to be represented as data structures instead of executable code. A lambda expression that is convertible to a delegate type `D` is also convertible to an expression tree of type `System.Query.Expression<D>`. Whereas the conversion of a lambda expression to a delegate type causes executable code to be generated and referenced by a delegate, conversion to an expression tree type causes code that creates an expression tree instance to be emitted. Expression trees are efficient in-memory data representations of lambda expressions and make the structure of the expression transparent and explicit.

The following example represents a lambda expression both as executable code and as an expression tree. Because a conversion exists to `Func<int, int>`, a conversion also exists to `Expression<Func<int, int>>`.

```
Func<int,int> f = x => x + 1;           // Code
Expression<Func<int,int>> e = x => x + 1; // Data
```

Following these assignments, the delegate `f` references a method that returns `x + 1`, and the expression tree `e` references a data structure that describes the expression `x + 1`.

Note

The structure of expression trees will be covered in a separate specification. This specification is not available for the May 2006 Technology Preview.